Routing in Fat Trees (RIFT) Hackathon

Chaos Monkey Testing of RIFT Implementations

IETF 104, Prague

Bruno Rijsman, 24-Mar-2019

What is RIFT?

- RIFT = Routing In Fat Trees
- A new link-state routing protocol optimized for “fat tree” topologies
- Main use case is large data center networks

RIFT Hackathon Participants

Artur Makutunowicz	LinkedIn
Bruno Rijsman	Individual
Pascal Thubert	Cisco
Tony Przygienda	Juniper Networks

RIFT Hackathon Activities

- Use open source RIFT-Python implementation.
- Generate configuration for large data center topology.
- Run topology in virtual machine (AWS instance).
- Automatically test correct initial convergence.
- Introduce a sequence of random "perturbations" in the topology.
- Automatically test correct initial re-convergence.
- For more details see:
  - http://bit.ly/rift-hackathon-ietf-104: hackathon instructions document
  - https://youtu.be/GqebgPmA4Xc: hackathon instructions video

Run fat tree topology in virtual machine

- Each RIFT-Python router runs in separate network namespace
- Use virtual Ethernet (veth) pairs to connect routers.

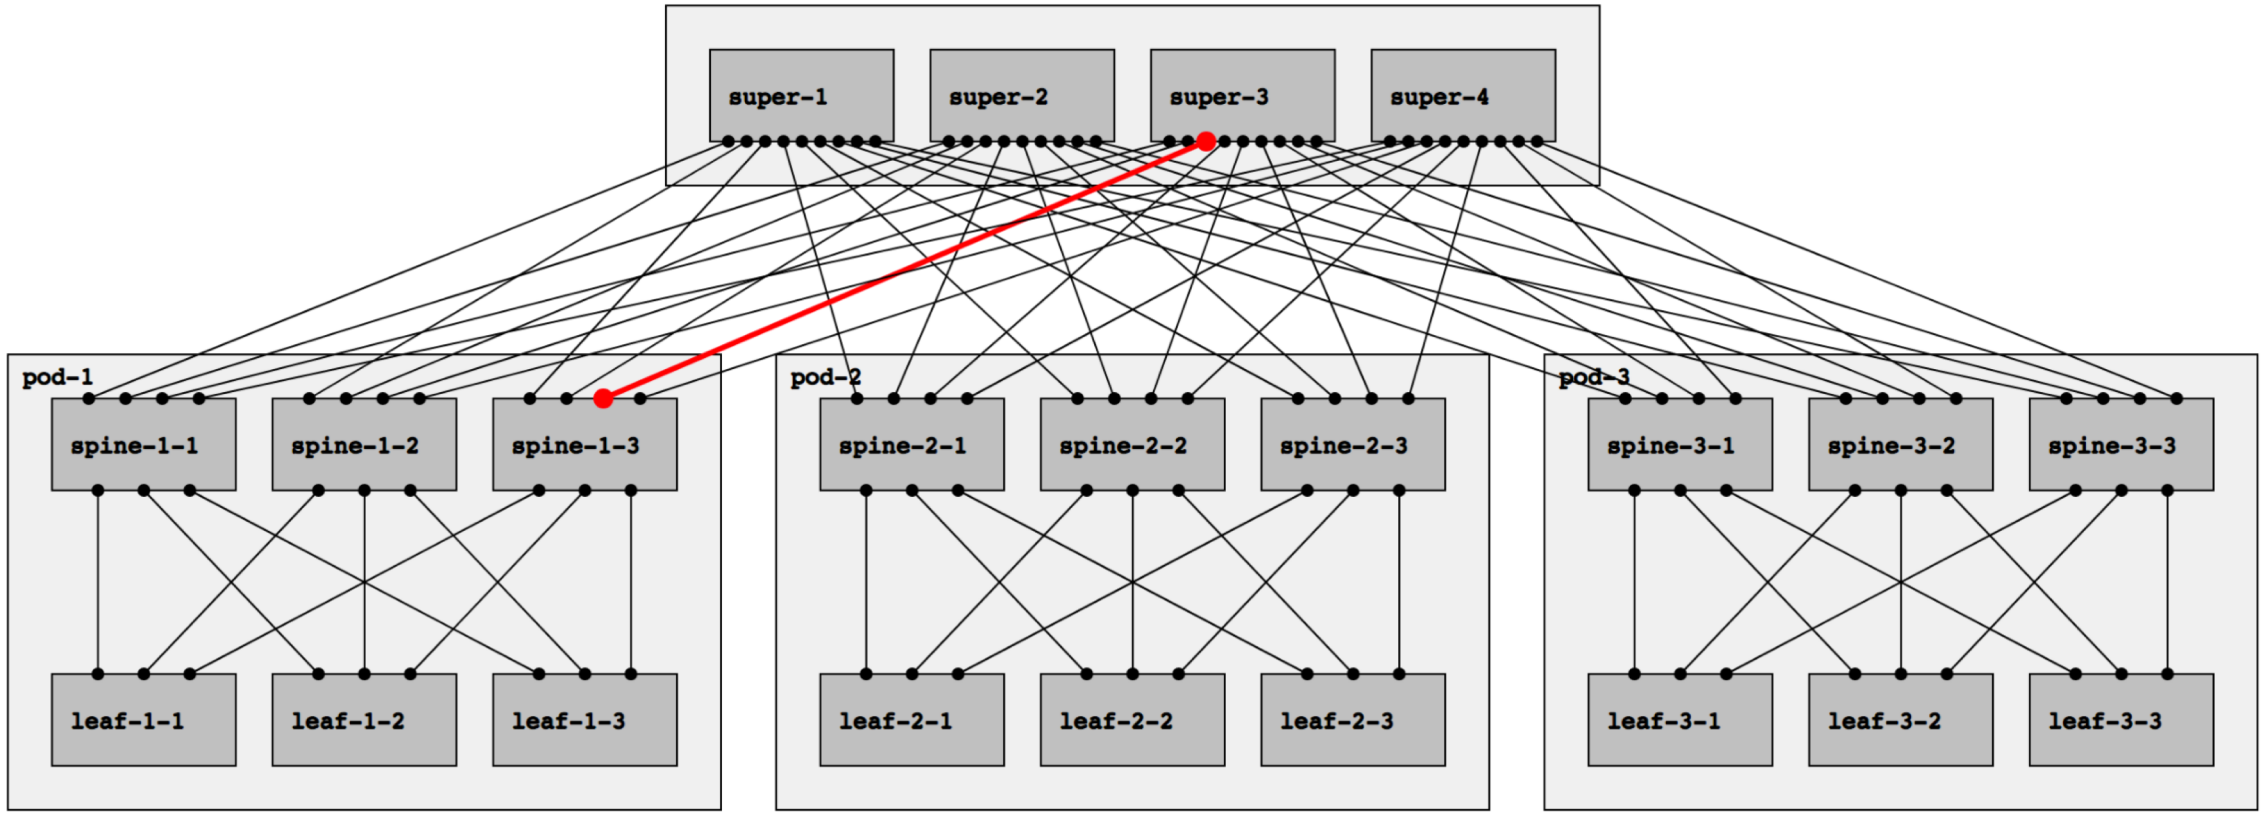

Generate topology and scripts

Meta-configuration

```
nr-pods: 3
nr-leaf-nodes-per-pod: 3
nr-spine-nodes-per-pod: 3
nr-superspine-nodes: 4
```

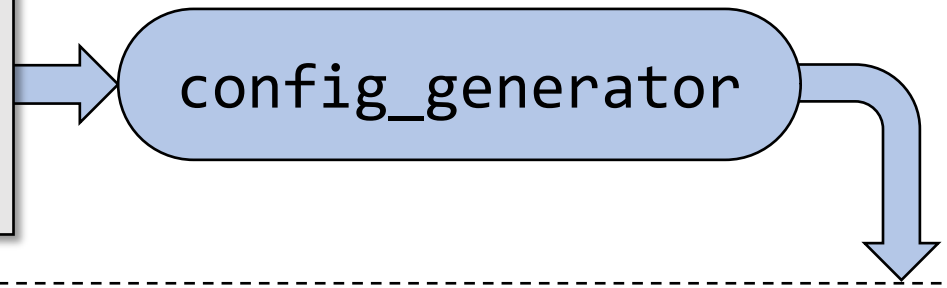

Configuration for each RIFT router

`leaf-1-1` `spine-1-1` `super-1` `…`

Scripts to start and stop topology

`start.sh` `stop.sh`

Scripts for “chaos monkey” perturbations

`chaos.sh`

Scripts to test correct convergence

`check.sh`

"Chaos monkey" perturbation testing

- Tool generates "chaos script" to randomly break and repair things
- "Chaos script" is topology aware
- Things that are being broken and repaired:
  - Full bi-directional link failures
  - Node failures
  - More things in future: uni-directional link failures, packets drops, packet re-ordering, packet delay, slow CPU, …
- All breakages are repaired at the end of the script
- After chaos script is finished, run check script to check convergence

"Chaos monkey" perturbation testing

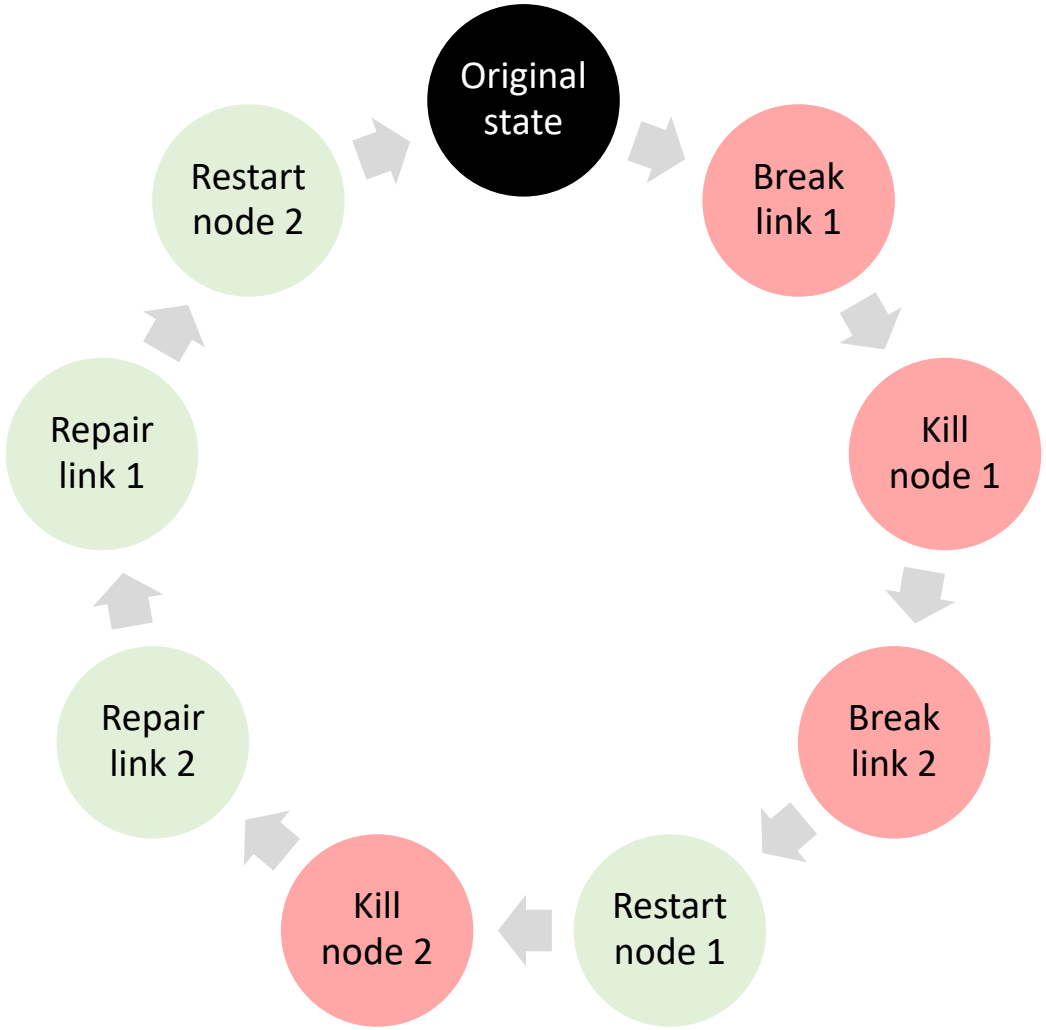

Automated convergence testing

- Tool generates “check script” to check correct re-convergence
- “Check script” is topology aware
- Things that are tested:
  - Ping from every leaf to every other leaf
  - Each node is up
  - All adjacencies are up (3-way)
  - North-bound default routes are in RIB and FIB and kernel
  - South-bound specific /32 routes are in RIB and FIB and kernel
  - More things in future

Protocol visualization tool to help debug issues

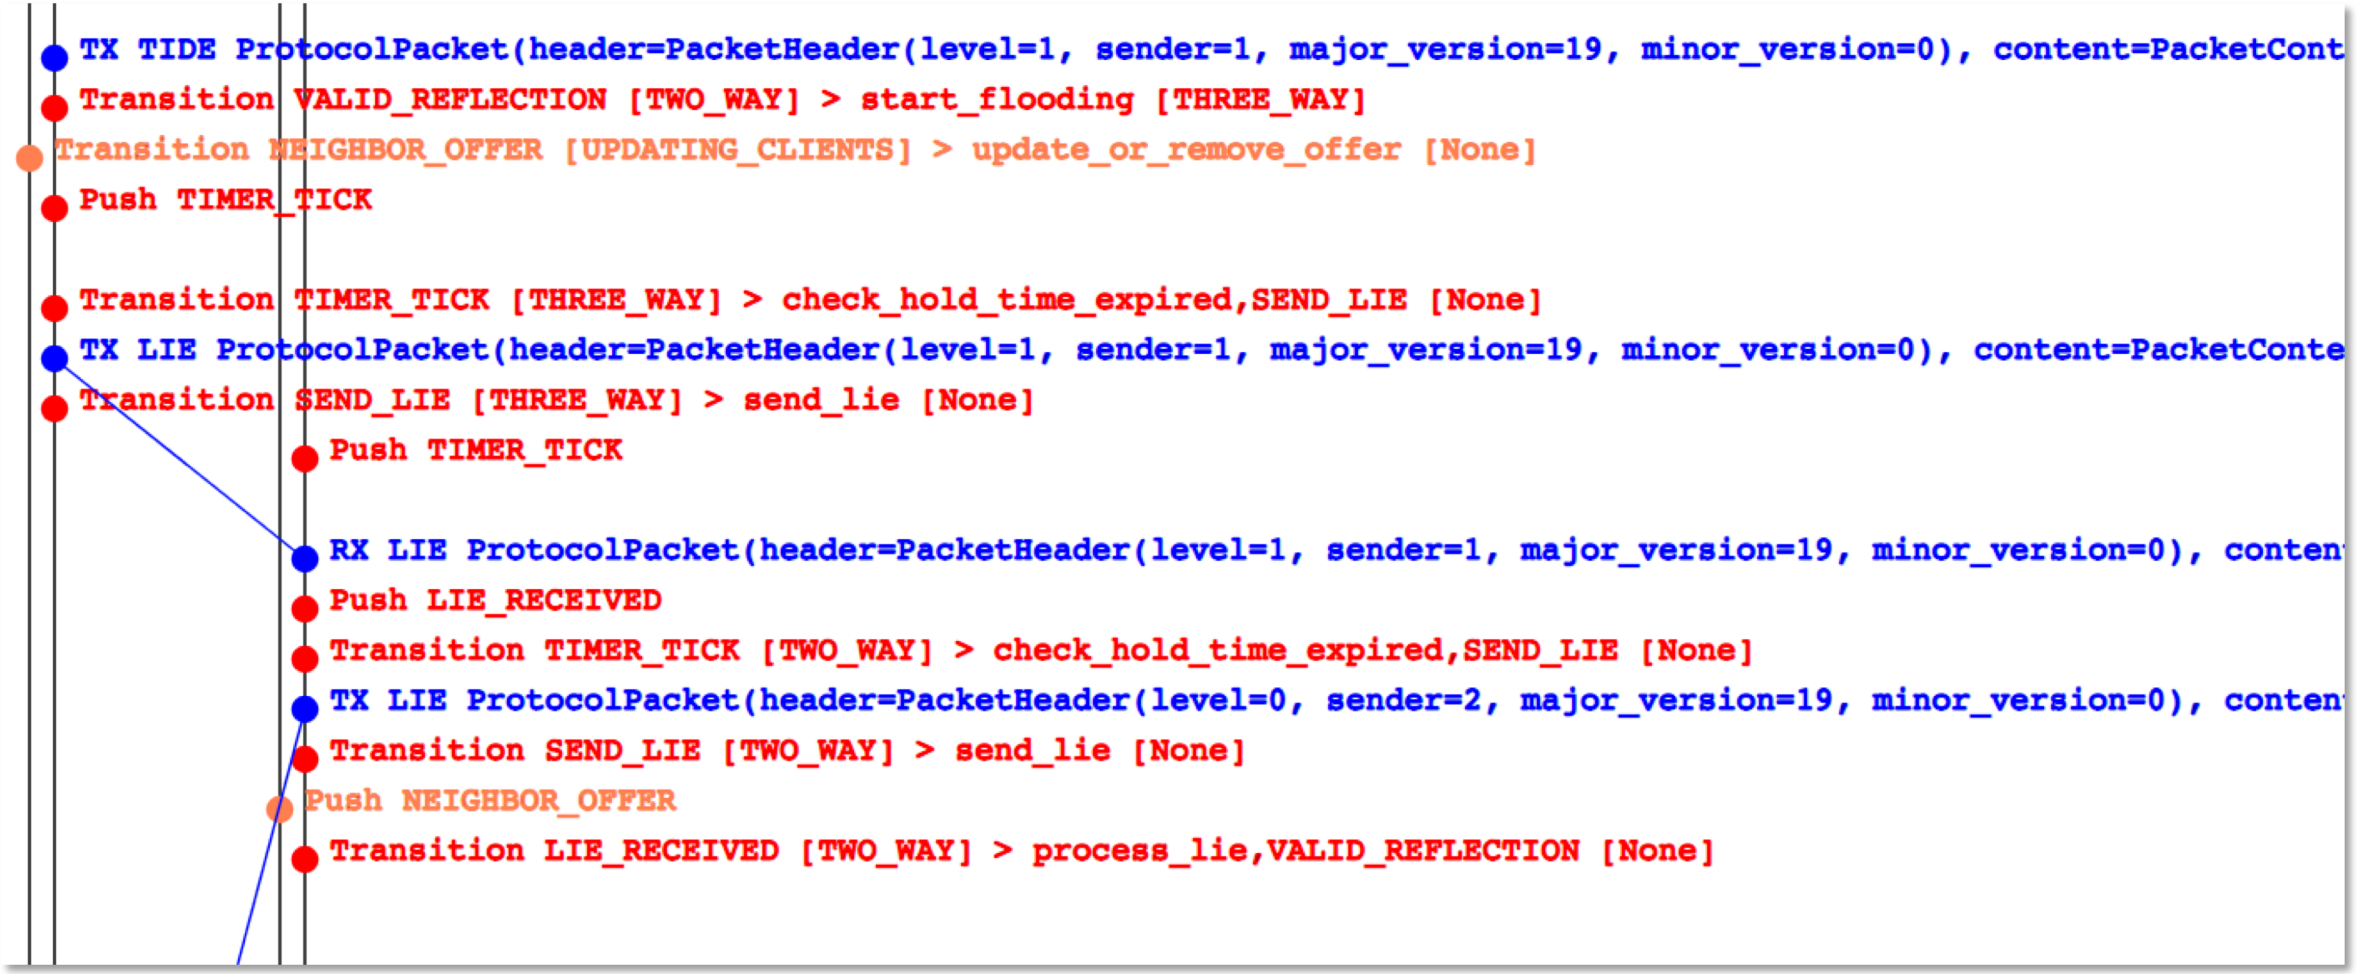

Hackathon results and lessons learned

- Implemented framework for:
  - Generate large fat tree topology (config files, scripts, visualization, …)
  - Run topology in virtual machine (AWS instance, using network namespaces)
  - Automated “chaos monkey” perturbation testing
  - Automated testing of correct re-convergence
- Lessons learned
  - We found and fixed several implementation issues using the framework:
  - IPv6 flooding issue (IPv4 in one direction, IPv6 in the other direction)
  - Multiple scenarios where exceptions are not handled in shut-down scenarios
  - Several ideas for new show commands to help debug issues (some implemented)